

Where it can get better — while staying grounded

Every option below pulls from real data (the manuals' own tables, your AMDF extract, your logs). None invent steps, specs, or stock numbers.

1 · Structured RPSTL parts

HIGH IMPACT

Parse the parts catalogs (RPSTL) into a real parts table: NSN ↔ part# ↔ FIG ↔ nomenclature ↔ vehicle. Adding a part auto-fills those fields exactly (cited), and powers look-alike variant warnings (your 'identical but different' goal).

Grounded: the manual's own parts tables · Effort: med-high · Helped by OCR

2 · FEDLOG / AMDF import

HIGH IMPACT

Import a FEDLOG/AMDF CSV the unit already has; auto-fill the FEDLOG row (unit price · AAC · ARC) by NSN. Today those three are typed by hand.

Grounded: your authoritative AMDF extract · Effort: low-med · Offline

3 · Verbatim procedure panel

VISION (goal C)

For a fault/part, surface the maintenance task verbatim & cited: remove/install steps, tools, torque — never paraphrased. Uses the unused 'procedures' table.

Grounded: verbatim WP text, cited · Effort: high · OCR-dependent

4 · Finish OCR + coverage meter

UNLOCKS ALL

Complete the GPU OCR pass and show a per-vehicle coverage meter ('72% searchable'). Sharpens search, schematics, and PMCS tech-status everywhere.

Grounded: makes scanned pages citable · Effort: med (CUDA runtime)

5 · Multi-sheet pagination

QUICK WIN

When a request exceeds 6 items, paginate cleanly across multiple 104th sheets instead of capping at 6.

Grounded: mechanical · Effort: low

6 · Capture confirm/override

QUICK WIN

Log when a mechanic accepts or overrides a tech-status (or part) suggestion, so the history signal and ranking get measurably better over time.

Grounded: your own confirmations · Effort: low

My recommendation: start with #2 (FEDLOG import) + #5/#6 (quick wins) for immediate end-to-end completeness, then #1 (structured RPSTL parts) as the big grounded accuracy leap, with #4 (OCR) running in parallel to unlock #1 and #3.