


More NSNs (append-only, R6) + the space/scope of growing the dataset

v0.21.0 · NSN enrichment now keeps every version forever (R6: add, never remove). And: what it costs to add NSNs vs whole TMs.

NEW RULE R6 — append-only: you may always ADD to the search engine, but never take away, even if the information is outdated.

APPEND-ONLY NSN ENRICHMENT (built)



Verified: two passes on one NSN → both versions retained; current = latest.

SPACE & SCOPE — adding NSNs vs TMs

NSNs for everything in your set

TEXT

tiny — ~0.5 KB/row; even 100k NSNs ≈ tens of MB. (GSA bulk CSV download is a few hundred MB, but you keep only your matches.)

Hardware / thread standards

TEXT

trivial — kilobytes (FED-STD-H28 dims).

More TM documents (full manuals)

BINARY

heavy — PDFs MB-tens of MB each + OCR/index; hundreds of TMs = GB-tens of GB + OCR time. Different sourcing (public TM repositories, not data.gov).

WHAT I'D DO

- Your definite want — 'new NSNs for everything' — is the cheap one. Run ``enrich --gsa <official GSA extract.csv>``: it appends current GSA name/desc/price for every NSN already in your index, keeping all prior versions (R6). Tens of MB at most.
- Standards/hardware: already seeded; trivial to extend.
- Adding whole TMs is the space-heavy, sourcing-sensitive part — best done deliberately by platform/manual, with version/authority checks. It's a different pipeline (crawl + OCR + index) than the NSN text enrichment.
- One practical note: the GSA Extract is a bulk download (hundreds of MB) — that's a file download, so I'll state size/source and get your OK, or you drop the CSV in and point ``--gsa`` at it.