


BUILT — Add documents: index new TMs without the command line (v0.50.0)

Fulfils the mission's 'any additional files added without a sweat' goal: point the program at a folder of PDFs and it indexes the new ones — safely.

THE FLOW (/ingest)



Already-indexed files are skipped (dedup by path + fingerprint); the crawl is resumable, so a big folder can be done in passes. Keep-alive makes the 2-second status polling cheap.

WHAT MAKES IT SAFE

- **Additive only**
the crawl pipeline only ADDS documents and text — it never deletes or overwrites your existing corpus or index (R6).
- **Snapshot first**
a safeguard snapshot is taken before any write, so the whole add is rollbackable (R1).
- **Reuses the tested path**
runs the same viewer_ingest.py 'crawl' that built the corpus — no new, unproven write code; progress comes from the runs table.
- **Read-only preview**
you see exactly what WOULD be added before committing — nothing is written during preview.

VERIFIED

`ingest_preview` validated on a synthetic folder: 4 PDFs across sub-folders, 1 already indexed → 3 new, bad path rejected, status idle when nothing runs. The write path shells out to the proven crawl CLI in the background. After adding, run the GPU OCR to make scanned pages searchable. New `/ingest` + 3 endpoints; rollback = remove them. JS lints clean.