

BUILT — Speed & efficiency pass (v0.56.0)

Five internal optimizations — no new UI, no new surface area. Faster paging, faster lookups, fewer bytes on the wire. All RPS-safe and additive.

TIER 1 — the big wins

☐ Open-PDF LRU cache

- `render_page_png` re-opened + re-parsed the PDF on EVERY page. Now an LRU of 8 open `fitz.Documents` is reused.
- PyMuPDF isn't thread-safe -> each doc carries a render lock; the highlight path still uses a fresh doc (it mutates).
- Big win for paging + the loupe.

Verified: LRU reuse + evict + close.

☐ Thread-local DB connections

- `db()` opened a new SQLite connection + re-applied PRAGMAs on every request (the dossier fires 6+).
- Now a per-thread connection is reused; callers' `.close()` is a harmless no-op; rebuilds if it goes bad.
- Relaxed/OCR mode still gets a fresh conn (no shared locks).

Verified: reuse, Row-factory, 5-thread x20, 0 errors.

☐ Indexed look-alike (NOCASE)

- Look-Alike scanned the whole parts table via `UPPER(name)=UPPER(?)`.
- Now name/nomenclature are matched with `COLLATE NOCASE` so the new NOCASE indexes apply.
- EXPLAIN: MULTI-INDEX OR across `ix_parts_name` + `ix_parts_nomenclature` on real (selective) data.

Index built by `optimize_index.py`.

TIER 2 — cheap & worth it

☐ ETag / 304 Not-Modified

- `_send` now sends an ETag (md5 of the body, stable across gzip). A matching If-None-Match returns 304 with an empty body.
- Repeat views of a page image, the JS, or JSON skip re-sending the bytes entirely — big on slow links.

Verified by curl: 304 on match, 200 on stale, Cache-Control preserved.

☐ Index maintenance + ANALYZE

- `optimize_index.py` (+ `.bat`): idempotent `CREATE INDEX IF NOT EXISTS` for `pages(document_id)` and `parts(name/nomenclature NOCASE)`, then `ANALYZE` so the planner uses them.
- EXPLAIN confirms find-in-manual now SEARCHes via `ix_pages_document` instead of scanning the biggest table.

Run once when OCR is paused (brief write lock; 120s busy-timeout).

SAFE, MEASURED, NO BLOAT

Every change is server-internal — no new pages, no API surface, nothing for the user to learn. All RPS-safe (pure Python, ES-agnostic). 44/44 tests still green; the connection-reuse `db()` is exercised by the passing pillar suite. The two index builds are the only DB writes and are a one-time, idempotent maintenance step — kept out of the live server so they never contend with the running OCR. This is efficiency by tightening hot paths, not by adding machinery.