

BUILT — OCR scan COMPLETE: the full corpus wired in (v0.66.0)

The text layer has been live in search/find/collections/callouts/3D as it filled. FINALIZE-OCR.bat (host-side) locks in the now-COMPLETE scan.

ALREADY LIVE (auto-wired by the FTS trigger as OCR filled each page):

Search · Find-in-manual · Smart Collections (auto-fill + 'new' badges) · page/schematic callouts · 3D part references. Every page became searchable the moment it was OCR'd — no rebuild.

1 Refresh structured parts

viewer_ingest.py parts re-extracts NSN/figure/nomenclature from every now-readable RPSTL page.

Parts index reflects the **WHOLE** scan.

2 Optimize the index

optimize_index.py rebuilds suggest_terms from the COMPLETE vocab + ANALYZE + WAL + indexes.

Predictive type-ahead now covers everything.

3 Milestone backup

safeguard snapshot --with-db takes a consistent copy of the finished index (+ off-disk mirror).

The finished corpus is protected.

4 Completion report

ocr_report.py --full writes the OCR completion report (pages, coverage, engine).

A record of the finished run.

5 Top nomenclatures

top_nomenclature.py answers 'which part comes up most' (battery / a specific bolt / gasket ...).

Answers the standing question #37.

6 Health check

verify_all + RPS lint confirm the additives pass and legacy parity holds after finalizing.

Green before and after.

ONE BUTTON, HOST-SIDE: FINALIZE-OCR.bat

Runs the six steps in order on Windows (the multi-GB index can't be written from a sandbox). top_nomenclature.py verified against the fixture; FINALIZE-OCR.bat + the scripts confirmed on host. After it, the complete scan is fully leveraged across the whole program.