

BUILT — most-common part nomenclature (v0.72.1)

Now that the whole corpus is OCR'd, rank every cited part and answer: which one comes up THE MOST — by name, by exact NSN, and by official FLIS name.

Three rankings

- Headline: most-common NOMENCLATURE — the cited RPSTL item / figure name (BOLT, GASKET, ...).
- Most-common EXACT NSN — the single physical part that recurs the most, with vehicle spread.
- Optional --flis: the official FLIS item name (INC -> H6) for the top NSNs, joined from PUB LOG.

Name, exact part, and official name.

One double-click

- ANSWER-MOST-COMMON-PART.bat runs it on the live index, prints the answer, and opens the result.
- Saves index\MOST-COMMON-PART.txt — copy the >>> ANSWER line straight back.
- Read-only open (mode=ro + query_only; immutable fallback) — safe to run while OCR / the app is live.

No command line needed.

Why on the PC

- The live index is 3.65 GB (891,556 x 4 KB pages per its own header).
- The dev sandbox mount serves it ~14.5 MB short (~3,500 pages) — SQLite reads that as 'malformed'.
- So this single number is computed on the machine that holds the whole file; the file itself is fine.

Truncating mount can't, your PC can.

VERIFIED (logic) + host run pending

Ranking SQL fixture-tested in isolation: headline nomenclature ('BOLT, MACHINE'), per-NSN frequency excluding nulls, COALESCE name fallback, and distinct vehicle/NSN counts all correct. The number itself comes from your run — double-click the bat; it reads index\viewer.db host-side (where it is coherent) and writes index\MOST-COMMON-PART.txt.

Additive (R1/R6): a new bat + an enhanced read-only script — no schema/route/corpus changes; the old `python top\_nomenclature.py` call still works. Optional FLIS join is name-only and read-only against PUB LOG CSVs.