


OCR — status & resume v0.99.16

The scan stalled at 43.8% on June 3. ~96% of the whole library is already searchable; the gap is the remaining scanned pages.

SCANNED PAGES NEEDING OCR — 121,135 total

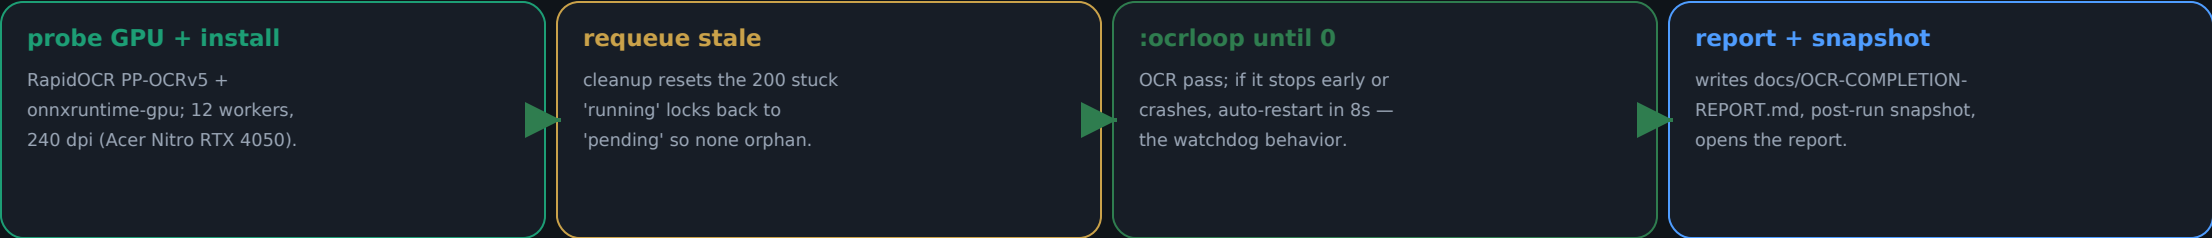


THE WHOLE LIBRARY — ~1.85 M pages across 39,683 documents



1,727,330 native-text + 53,016 OCR'd = 1,780,346 searchable now (96.3%). Remaining 68,119 scanned pages → OCR.

RESUME-OCR.bat → run to 100% (self-restarting)



MONITOR — daily task 'viewer-ocr-daily-progress' (09:08)

Reads ocr_progress_history.tsv (the runner appends to it live) + the heartbeat freshness — NOT a slow full-table scan — and pings Chris each morning with % done, pages remaining, running-vs-stalled, and a reminder to launch RESUME-OCR.bat if it stopped. At 100% it says the library is fully OCR'd and the check can be turned off.

TO START: double-click RESUME-OCR.bat (plug in AC, fans high; ~15-20h GPU time, resumable). Desktop automation can't start the long run for you.