

v1.13.0 Holistic Review — implement · audit · harden · polish

3 parallel reviewers (arch/perf · robustness/security · data-integrity/UX) → recommendations → built & re-verified green

Architecture / Perf

features/corpus.py: 1 leak-free
FTS retrieval (was ~15 leaky connects)
doc_path() pooled lookup
auto-optimize: WAL now +
bg index build (OCR-safe)
bounded worker pool
fast state = shipped default,
not a manual step

Robustness / Security

_sent flag: no double-send
(B9 boundary now structural)
all sidecar writes → atomic_write
(fsync + _replace_retry)
chunked-TE reject · temp-PDF
try/finally · layout/tables 500→empty
exposure token off-loopback
durable-write audit guard added

Data-integrity / UX

/measures validates + WITHHOLDS
quarantined garble (R13)
/part dims: real page cite or
honest "verify" — never faked
safety callouts cite RESTORED
(refactor had silently dropped them)
home-page kiosk/a11y fixed
Tools menu: +part/measures/scan

AUDIT of the new work (handle with care)

2 of 3 review subagents WROTE code despite read-only → treated as unverified, re-ran the full battery.
Caught 1 real regression: corpus key rename cautions→results dropped SAFETY callouts from /part. Fixed.

GREEN after audit

11 regression suites · 279 route cases · 24k fuzz · audit 0 FAIL/0 WARN · durable-write guard PASS
additive + rollback-safe (R1) · VERSION 1.13.0 · VERIFY.bat = one exit-code-truth gate

"A passing suite that never runs the code is worse than no suite." — the audit exists to make that impossible.